

# The CAP Theorem

## Availability, Consistency, Failure in Distributed Systems

### Distributed Systems

Andrea Omicini

<mailto:andrea.omicini@unibo.it>

Dipartimento di Informatica – Scienza e Ingegneria (DISI)

ALMA MATER STUDIORUM – Università di Bologna

Academic Year 2025/2026



- 1 Availability, Consistency, Failure
- 2 CAP at a Glance
- 3 A Brief Detour: Pokémon Go
- 4 Back to CAP: Some Proof
- 5 Building on CAP
- 6 Conclusion



# Next in Line...

- 1 Availability, Consistency, Failure
- 2 CAP at a Glance
- 3 A Brief Detour: Pokémon Go
- 4 Back to CAP: Some Proof
- 5 Building on CAP
- 6 Conclusion



# (Hiding) Failure in Distributed Systems

- being able to **keep on providing services in spite of failures** is supposed to be one of the main benefit of distributed systems over centralised ones <sup>[Friedman and Birman, 1996]</sup>
- intuitive assumption: **distributed systems can be designed** so that if one component of the system fails – or, it becomes disconnected / partitioned – other components can replace it so as **to hide failures** from the outside world
  - or at least to reduce the (perceived) impact of failure



# Failure & Availability

- when a system can hide (most of) its failures, it is basically (almost) *always working*
  - we say it is *highly available*
- basic questions in the design of distributed systems: *how much failure can a given system sustain before failure is noticed?*



# Features

What do we expect from a (distributed) systems?<sup>[Gilbert and Lynch, 2012]</sup>

- we would like them to behave *correctly*, to give correct replies when queried
  - to be **consistent**
- we would like them to work, to make good things happen, to be *live*
  - to be **available**
- however, we know systems *can* actually experience power losses, crashes, network failures, message loss, malicious attacks, Byzantine failures, and so on
  - and be **unreliable**

# Tradeoff

## The idea

- “The CAP theorem is one example of a more general tradeoff between *safety* and *liveness* in *unreliable* systems” [Gilbert and Lynch, 2012]




# Next in Line...

- 1 Availability, Consistency, Failure
- 2 CAP at a Glance**
- 3 A Brief Detour: Pokémon Go
- 4 Back to CAP: Some Proof
- 5 Building on CAP
- 6 Conclusion





# Informally I



© Mr. Lovenstein

# Informally II

In short<sup>[Shim, 2012]</sup>

According to the CAP theorem, it is only possible to simultaneously provide any two of the three following properties in distributed applications: **consistency** (C), **availability** (A), and **partition tolerance** (P).



# Informally III

Less short<sup>[Zhao, 2014]</sup>

The CAP theorem [...] states that it is impossible to satisfy all three of the following guarantees:

**Consistency** (C) the replicated data is always consistent with each other

**Availability** (A) the data is highly available to the users

**Partition tolerance** (P): the system can continue providing services to its users even when the network partitions



# Original Theorem

First formulation by Eric A. Brewer<sup>[Brewer, 2000]</sup>

A distributed database potentially features three desirable properties

- 1 Consistency
- 2 Availability
- 3 tolerance towards network *Partition*

According to the CAP theorem, any shared-data distributed system can have at most two of these three properties

<http://apice.unibo.it/xwiki/bin/view/Talk/BrewerPodc2000>

# An Impossibility Result

- “we repent and renounce”
- ? ... yet, what is actually our sin?
- **distributed systems**, that's what it is.



# Pick Up One... I

... and live with(out) that

- one might *forget about tolerance to network partition*, so to have consistency and availability
- one might *forget about consistency*, and live with a partition-tolerant and highly-available system
- one might *forget about availability*, and enjoy a system that is both consistent and tolerant to network partition

# Pick Up One... II

## Note

- the three properties are not exactly of the same sort, both technically and conceptually
    - consistency and availability range over a spectrum of options, whereas partition tolerance can somehow be seen more as an on/off feature
  - all of them are *desirable*, yet forfeiting partition tolerance is *not really an option* in real-world systems
  - particularly with pervasive systems in the IoT era, where *instability* (of the network) rules<sup>[Grimm et al., 2004]</sup>
- long story short, CAP theorem most often forces us to *choose between availability and consistency*

# Example: Location-based Games I

## Location-based games

- *location-based services*<sup>[Shekhar et al., 2016]</sup> use information about the position of a user / device to provide information, entertainment, security
- *location-based games* use player's location to evolve and progress the gameplay
  - e.g., BotFighters, GeoZombie<sup>[Prandi et al., 2016]</sup>, Ingress, Pokémon Go, Harry Potter: Wizards Unite, Minecraft Earth



## Example: Location-based Games II

### CAP in location-based games: partition tolerance

- since the architecture is built around the notion of millions of players roaming physical space *worldwide* with their mobile devices (and, playing through them), forfeiting **partition tolerance** is *not* an option
  - mobile devices are inherently *unstable* in their network connectivity
  - players are inherently *mobile*, and so they can move in and out of network coverage
  - players tend to concentrate in some areas, e.g., during special events
  - scalability is a multi-level issue: not just the number of players overall, but also the number of places they can be at, and the number of players at each place
    - e.g., costly disaster at the Chicago 2017 Pokémon Go Fest

## Example: Location-based Games III

### CAP in location-based games: consistency and availability

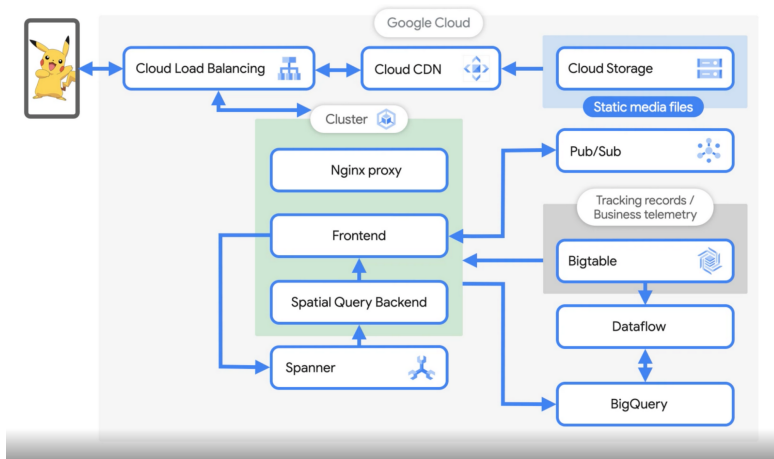
- **consistency** of the game data (goals, situation, achievements, ...) is essential to keep players going (and, into the game)
  - location-based games are usually built around logically-centralised architecture
  - using spatial replication in order to reduce user-perceived latency and improve scalability
  - when strong consistency is required, e.g., for in-game transactions, players might be forced to wait for the app servers to confirm the operation
    - by reaching a consistent state overall
- first (and usual) casualty: **availability**
  - the Spinning Wheel of Death 🎡

# Next in Line...

- 1 Availability, Consistency, Failure
- 2 CAP at a Glance
- 3 A Brief Detour: Pokémon Go**
- 4 Back to CAP: Some Proof
- 5 Building on CAP
- 6 Conclusion



# Catching a Pokémon on Google Cloud I



<https://cloud.google.com/blog/topics/developers-practitioners/how-pokémon-go-scales-millions-requests>

# Catching a Pokémon on Google Cloud II

- ❶ when a player opens the Pokémon Go app, all static media are downloaded to his/her personal device
  - which are stored in Cloud Storage
  - **Cloud CDN (Content Network Delivery)** uses Google's global *edge network* to serve content closer to users
  - Cloud CDN works with the **Cloud Load Balancing** (actually, a load balancer) to *cache* and serve user content
- ❷ user requests are sent to **NGINX** reverse proxy
  - which sends this traffic to the Frontend game service
  - hosted on **Google Kubernetes Engine (GKE)**
- ❸ the Spatial Query Backend handle the location-based features
  - keeps a cache *sharded* by location
  - the cache and service decides which Pokémon is shown on the map, what gyms and PokéStops are around, which is the time zone, ...

# Catching a Pokémon on Google Cloud III

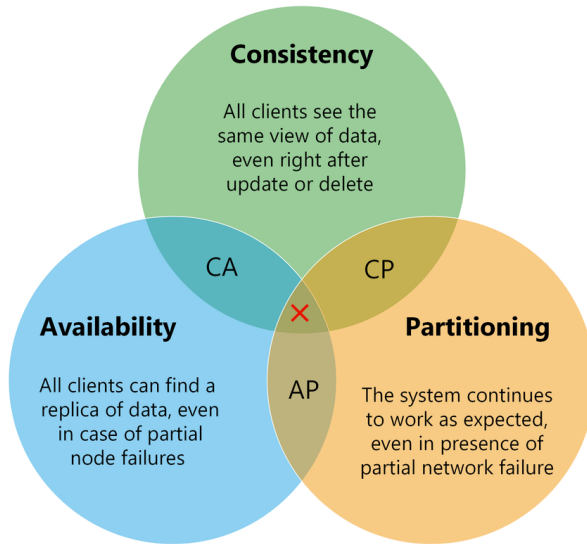
- 4 when a player catches a Pokémon, the Frontend (GKE) sends an event to **Google Spanner**
  - when the write request is completed, a response is sent back to the device
  - Spanner is *strongly consistent*
- 5 every player's actions is recorded in the **Bigtable** NoSQL database service
  - as a **Protobuf** representation
  - for logging and tracking data
  - also as a message to a **Pub/Sub** topic for the analysis pipeline
- 6 multiple players in the same geographical region are kept in sync by *determinism* of the Pokémon Go maps
  - even if multiple players are on different machines, but in the same physical location, all players' inputs would be the same
- 7 *all the servers are in sync with settings changes and event timings in order for all players to feel like they are part of a shared world*

# Next in Line...

- 1 Availability, Consistency, Failure
- 2 CAP at a Glance
- 3 A Brief Detour: Pokémon Go
- 4 Back to CAP: Some Proof**
- 5 Building on CAP
- 6 Conclusion



# One Way to See CAP





# From Brewer's Conjecture to CAP Theorem

## Going formal

- Brewer's first results required some *proof*, possibly formal
- for a proof, a more precise formulation of the conjecture is required
- a proof was soon provided<sup>[Gilbert and Lynch, 2002]</sup>
  - with some limitations
- then, several others followed
- first step for us to sketch a proof: defining the main concepts

# Consistency, Availability, Network Partition I

## Availability<sup>[Gilbert and Lynch, 2002]</sup>

- “For a distributed system to be continuously **available**, every request received by a non-failing node in the system must result in a response”
- if the system is available, we got *responses*



# Consistency, Availability, Network Partition II

## Consistency<sup>[Gilbert and Lynch, 2002]</sup>

- a **consistent** service is modelled as an *atomic data object*
- where
  - operations are totally ordered, and
  - each operation occurs in a single instant of time
- ! here, the meaning of *consistent* is unlike ACID properties, since it encompasses both A and C there
- among the many consequences, consistency implies that all *read* operations over a distributed shared memory occurring after a *write* operation completes must return the value of either this write operation or a later one
- if the system is consistent, we got *correct responses*

# Consistency, Availability, Network Partition III

## Network partition<sup>[Gilbert and Lynch, 2002]</sup>

- “When a network is partitioned, all messages sent from nodes in one component of the partition to nodes in another component are lost”
- “And, any pattern of message loss can be modelled as a temporary partition separating the communicating nodes at the exact instant the message is lost”

## In short, and roughly. . .

- if node A send a message to node B, the network is *partitioned* if the message does not make it to B
- *availability* is when B receives the message and responds
- *consistency* is when B response is correct

# Theorem I

## Three different sorts of systems

The proof is given for three sorts of network

- asynchronous network with message loss
- asynchronous network without message loss
- partially synchronous network with local clocks

For teaching purposes, we focus here on the *asynchronous model*, where

- there is no single clock
- nodes act based on local computation and message received

# Theorem II

## Theorem<sup>[Gilbert and Lynch, 2002]</sup>

“It is impossible in the asynchronous network model to implement a read/write data object that guarantees the following properties:

- availability
- atomic consistency

in all fair executions (including those in which messages are lost)”



Proof<sup>[Gilbert and Lynch, 2002]</sup> |

## Assumptions

- atomicity, availability, and partition tolerance are all fulfilled
  - *proof by contradiction* attacks just that
- nodes in the network can be partitioned into two disjoint, non-empty sets  $G_1, G_2$
- atomic object  $o$  has initial value  $v_0$ 
  - $o$  is expected to be consistent through  $G_1, G_2$
- $\alpha_1$  executes a single write  $v_1 \neq v_0$  of  $o$  in  $G_1$ 
  - $\alpha_1$  is the only request in that time
  - during  $\alpha_1$  no message are received—from  $G_1$  to  $G_2$ , from  $G_2$  to  $G_1$
- no messages from  $G_1$  are received in  $G_2$
- $\alpha_2$  executes a single read of  $o$  in  $G_2$ 
  - during  $\alpha_2$  no message are received—from  $G_1$  to  $G_2$ , from  $G_2$  to  $G_1$

Proof<sup>[Gilbert and Lynch, 2002]</sup> II

Q.E.D.

- due to availability,  $\alpha_1$  completes ( $v_0 \rightarrow v_1$ ), and  $\alpha_2$  does, too
- executing  $\alpha_1$  and  $\alpha_2$ ,  $G_2$  just sees  $\alpha_2$ , and must return  $v_0$
- this obviously violates consistency as defined above
  - which is basically a sort of atomic consistency





# Next in Line...

- 1 Availability, Consistency, Failure
- 2 CAP at a Glance
- 3 A Brief Detour: Pokémon Go
- 4 Back to CAP: Some Proof
- 5 Building on CAP**
- 6 Conclusion



# What do we Make out of an Impossibility Result?

## Should we stop distributing computational systems?

- in the same way we stopped using axiomatic systems after Gödel?<sup>[Gödel, 1931]</sup>
  - of course we did *not*
- ! **negative results** just *set the boundaries*, and make us understand the very *limits of our reach*
- so that an impossibility result from computer science becomes a leverage for specific and effective computer engineering methods and practices

# Switching Strategy based on Partition Tolerance

## A very simple scheme

- Brewer himself later elaborated on his CAP theorem<sup>[Brewer, 2012]</sup>
  - (i) when the network is partitioned, a distributed system should choose a tradeoff between consistency and availability
  - (ii) when there is no partition, a system can feature both consistency and availability



# ACID vs. BASE I

## Beyond ACID<sup>[Fox et al., 1997]</sup>

- ACID (Atomicity, Consistency, Isolation, and Durability) semantics might be too strong
  - “The design space for network services can be partitioned according to the data semantics that each service demands. ”
  - “For other Internet services, however, the primary value to the user is not necessarily strong consistency or durability, but rather high availability of data”

# ACID vs. BASE II

## Towards BASE<sup>[Fox et al., 1997]</sup>

- “Stale data can be temporarily tolerated as long as all copies of data *eventually* reach *consistency* after a short time”
- “*Soft state*, which can be generated at the expense of additional computation or file I/O, is exploited to improve performance; data is not durable”
- “*Approximate answers* (based on stale data or incomplete soft state) delivered quickly may be more valuable than exact answers”

## BASE

- Basically Available
- Soft state
- Eventual consistency

# Beyond BASE

## Data consistency for Cloud<sup>[Birman et al., 2012]</sup>

- most cloud services nowadays adopts BASE
  - e.g. eBay, Amazon DynamoDB
    - ! Amazon S3 allows for *strong read-after-write* consistency
    - ! Amazon DynamoDB writes are eventually consistent, reads too—but strongly-consistent reads can be configured (and, they are more costly)
- and try to mask inconsistencies from users
- steps ahead are leading to new models, based on, e.g., new data-consistency models
  - ! further details in the literature—at the end of the slides

# Overall

## CAP today<sup>[Brewer, 2012]</sup>

*The modern CAP goal should be to maximise combinations of consistency and availability that make sense for the specific application. Such an approach incorporates plans for operation during a partition and for recovery afterward, thus helping designers think about CAP beyond its historically perceived limitations*



# Next in Line...

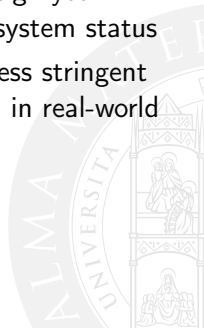
- 1 Availability, Consistency, Failure
- 2 CAP at a Glance
- 3 A Brief Detour: Pokémon Go
- 4 Back to CAP: Some Proof
- 5 Building on CAP
- 6 Conclusion**





# Lessons Learnt

- in the design and development of distributed systems, we have typically to choose between a responsive system and a fully-consistent one, when network fails
- in any case, we can leverage on the CAP theorem to design your system with different features depending on the overall system status
- ACID is generally a thing to know and understand, yet less stringent models could be both theoretically and practically useful in real-world distributed systems—e.g., BASE



# The CAP Theorem

## Availability, Consistency, Failure in Distributed Systems

### Distributed Systems

Andrea Omicini

<mailto:andrea.omicini@unibo.it>

Dipartimento di Informatica – Scienza e Ingegneria (DISI)

ALMA MATER STUDIORUM – Università di Bologna

Academic Year 2025/2026



# References I

[Birman et al., 2012] Birman, K., Freedman, D., Huang, Q., and Dowell, P. (2012).  
Overcoming CAP with consistent soft-state replication.  
*Computer*, 45(2):50–58  
DOI:10.1109/MC.2011.387

[Brewer, 2012] Brewer, E. (2012).  
CAP twelve years later: How the “rules” have changed.  
*Computer*, 45(2):23–29  
DOI:10.1109/MC.2012.37

[Brewer, 2000] Brewer, E. A. (2000).  
Towards robust distributed systems (abstract).  
In *19th Annual ACM Symposium on Principles of Distributed Computing (PODC '00)*, page 7, New York, New York, USA. ACM Press  
DOI:10.1145/343477.343502

[Fox et al., 1997] Fox, A., Gribble, S. D., Chawathe, Y., Brewer, E. A., and Gauthier, P. (1997).  
Cluster-based scalable network services.  
*ACM SIGOPS Operating Systems Review*, 31(5):78–91  
DOI:10.1145/269005.266662

[Friedman and Birman, 1996] Friedman, R. and Birman, K. (1996).  
Trading consistency for availability in distributed systems.  
Technical report, Cornell University  
(APICe) <https://hdl.handle.net/1813/7235>



# References II

- [Gilbert and Lynch, 2002] Gilbert, S. and Lynch, N. (2002).  
 Brewer's conjecture and the feasibility of consistent, available, partition-tolerant web services.  
*ACM SIGACT News*, 33(2):51  
 DOI:10.1145/564585.564601
- [Gilbert and Lynch, 2012] Gilbert, S. and Lynch, N. (2012).  
 Perspectives on the CAP theorem.  
*Computer*, 45(2):30–36  
 DOI:10.1109/MC.2011.389
- [Gödel, 1931] Gödel, K. (1931).  
 Über formal unentscheidbare sätze der Principia Mathematica und verwandter Systeme I.  
*Monatshefte für Mathematik und Physik*, 38(1):173–198  
 DOI:10.1007/BF01700692
- [Grimm et al., 2004] Grimm, R., Davis, J., Lemar, E., Macbeth, A., Swanson, S., Anderson, T., Bershad, B., Borriello, G., Gribble, S., and Wetherall, D. (2004).  
 System support for pervasive applications.  
*ACM Transactions on Computer Systems*, 22(4):421–486  
 DOI:10.1145/1035582.1035584
- [Prandi et al., 2016] Prandi, C., Roccetti, M., Salomoni, P., Nisi, V., and Nunes, N. J. (2016).  
 Fighting exclusion: a multimedia mobile app with zombies and maps as a medium for civic engagement and design.  
*Multimedia Tools and Applications*, pages 1–29  
 DOI:10.1007/s11042-016-3780-9

# References III

- [Shekhar et al., 2016] Shekhar, S., Feiner, S. K., and Aref, W. G. (2016).  
Spatial computing.  
*Communications of the ACM*, 59(1):72–81  
DOI:10.1145/2756547
- [Shim, 2012] Shim, S. S. (2012).  
Guest editor's introduction: The CAP theorem's growing impact.  
*Computer*, 45(2):21–22  
DOI:10.1109/MC.2012.54
- [Zhao, 2014] Zhao, W. (2014).  
*Building Dependable Distributed Systems*.  
Wiley  
DOI:10.1002/9781118912744

